

Context-Aware Natural Language Processing for Interpretable Analysis of Pharmaceutical Deviation Narratives

Author: Erin Xu

Manager: Saeed Babaeizadeh

Team: R&D Quality

Github repository: QDN space, TBA

Purpose

This document is intended for technical stakeholders to reference when reviewing the proof-of-concept implementation of a context-aware natural language processing (NLP) enrichment and retrieval workflow for pharmaceutical deviation records. It describes first the problem statement, objectives, data flow, and storage organization, inputs, outputs produced by the pipeline, and the technical dependencies required to run the workflow. It also explains how to use the semantic-search layer to build retrieval cohorts for specific business questions, and documents the hard-coded statistics, thresholds, and heuristics configured in the reporting layer—so that every analyst-set value can be audited and adjusted. For any unresolved curiosity, there is also a section outlining the justification of this methodology as well as the conceptual breakdowns of the scientific concepts behind this technology. Finally, future work and next steps are outlined.

Problem Statement

Quality Data Navigator (QDN) provides a governed platform for structured quality data, but the narrative text associated with pharmaceutical deviations remains difficult to query and analyze systematically. These narratives contain operational context, referenced entities, recurring themes, and root-cause signals that are not fully represented in existing structured fields.

Objectives

This project develops a context-aware NLP enrichment and retrieval workflow for pharmaceutical deviation records. Rather than analyzing only the deviation title and description, the workflow constructs tiered text representations across event narratives, quality and impact assessments, structured root-cause descriptions, and action text, at both event and action-record grains. These representations support entity extraction and normalization, contextual enrichment, embeddings, topic and root-cause analysis, vector indexing, and semantic and hybrid retrieval.

Success is measured by the ability to:

1. Demonstrate that heterogeneous free-text fields can be consolidated into reproducible core, mid, and full deviation representations while preserving event and action-record grain.
2. Extract and normalize relevant entity-level information from the available narrative and descriptive fields.
3. Connect deviation records with relevant contextual and reference sources.
4. Produce enriched text suitable for embeddings and downstream NLP analysis.
5. Support both concept-based and exact-term retrieval.
6. Identify interpretable topics and root-cause signals.
7. Retain links to source fields and supporting evidence for review.
8. Produce a supplementary dataset that can be joined back to QDN data using a documented event identifier.

The following capabilities are out of scope for this project:

- Production deployment, operational support, or service-level commitments.
- Direct integration of the supplementary dataset or retrieval indexes into the production QDN application. This project produces a supplementary dataset and retrieval indexes that can be joined back to QDN data using a documented event identifier, but does not modify the production QDN application or its governed data.
- Real-time or incremental processing of newly created or updated deviation records. Production serving, continuous processing, operational monitoring, retraining, and lifecycle management require separate productionization and governance work.
- Automated quality decisions, record disposition, compliance determinations, or replacement of subject-matter-expert review.
- Changes to source-system records or existing governed QDN fields.
- Automated remediation of missing, inconsistent, or incorrect source data.
- Comprehensive extraction of every possible entity type or terminology variant.

Suggestions of these practices will be outlined in the final section.

Overview of Data Flow

The pipeline combines formatted TrackWise deviation records with six supporting reference datasets. Reference records are normalized into a common lookup contract and consolidated into a unified glossary. Deviation records are then compressed to one row per event, cleaned, filtered to the supported language scope, resolved against the glossary, and represented at three levels of narrative detail.

A Databricks-hosted foundation model generates a short situating context for each deviation. The deterministic reference context, generated context, and source text are combined into contextual retrieval text. BGE-M3 converts each text tier into a 1,024-dimensional dense vector. Three Databricks AI Search indexes expose the resulting core, mid, and full representations for hybrid vector and keyword retrieval.

The two figures below show the big picture overview of the workflow at decreasing levels of abstraction. The pipeline resolves reference entities before generating situating context and constructing the three embedding representations.




Figure 1: Principal data flow of the context-aware NLP enrichment and retrieval pipeline.




Figure 2: Stakeholder view of the pipeline and its supplementary outputs.

Engineering: Text Preparation and Filtering

Before text construction, the pipeline:

1. Applies Unicode normalization.
2. Standardizes Unicode hyphens and non-breaking spaces.
3. Folds accented Latin characters for matching.
4. Normalizes repeated spaces and excessive blank lines.
5. Preserves letter case because acronym and CRO extraction depend on case-sensitive source information.
6. Applies language detection to the combined `Event_Title` and `Event_Description`.
7. Retains records classified as English, including supported code-only records.
8. Applies an additional guard against remaining CJK characters in the title and description.
9. Excludes records whose resulting full text is empty.

Language detection is seeded to make repeated runs reproducible.

The English-language filter is based on the core title and description, not independently on every source field. This behavior should be retained in the documented acceptance criteria and reviewed if multilingual coverage is later required.

Engineering: Event-Grain Reconciliation

The source table may contain multiple rows for the same `Event_Number` because `Action_Text` is handled at a row grain below the deviation event. Before downstream enrichment, the pipeline compresses the source to one row per event:

- Event-stable fields use the first available non-null value.
- Distinct non-null `Action_Text` values are collected and joined into a single text field.
- `Event_Number` is cast to a string and renamed `pr_id`.
- Duplicate action text is removed during aggregation.

This produces the event-level grain expected by the embedding table, vector indexes, and supplementary outputs.

Storage Organization

All persistent data objects used by the proof of concept are located in the development Unity Catalog.

Logical layer	Namespace	Contents
Development catalog	<code>us_gmsgq_dev</code>	Catalog containing the proof-of-concept data objects
Mart layer	<code>us_gmsgq_dev.gms_us_mart</code>	Read-only formatted deviation source
Analytics layer	<code>us_gmsgq_dev.gms_us_alyt</code>	Reference inputs, normalized lookups, enrichment tables, embeddings, and retrieval indexes

The mart table is treated as a read-only input. The pipeline writes derived data products to the analytics schema. These are development resources and do not represent a production deployment.

Input: Formatted TrackWise Deviation Dataset

Attribute	Value
Logical name	Formatted TrackWise deviation data
Fully qualified table	<code>us_gmsgq_dev.gms_us_mart.tw_deviation_data_formatted_rdq</code>
Pipeline role	Primary source of deviation identifiers, narrative fields, assessments, root-cause fields, and action text
Read/write behavior	Read only
Pipeline event identifier	<code>Event_Number</code>
Derived event identifier	<code>pr_id</code> , created by casting <code>Event_Number</code> to string
Target processing grain	One row per deviation event

Input: Reference Data

The pipeline uses six reference datasets to resolve entities found in deviation records. Sources that live in Databricks (the Veeva controlled-document table and the vendor MDM) refresh on their own and are marked Dynamic; the SharePoint and CTMS exports are point-in-time snapshots and are marked Static.

Entity domain	Source location	Re-fresh	Raw input table	Primary source fields used	Pipeline purpose
Clinical identifiers	Excel, SharePoint, self-refreshing MDM	Static	us_gmsgq_dev.gms_us_alyt.ClinicalId deviations	Parent_Node, Name, Protocol Number, Alternate_Name, Development_Name, Parent_Alias, Grand_Parent, Grandparent_Alias, and available contextual attributes	Resolve clinical program, compound, protocol, and study identifiers
Controlled documents	Veeva (Databricks)	Dy-namic	us_gmsgq_dev.gms_us_alyt.documents_number_ deviations	document_number__v, previous_document_number__c, legacy_numbers__c, title__v	Resolve current, previous, and legacy controlled-document numbers
Acronyms	SharePoint	Static	us_gmsgq_dev.gms_us_alyt.Acronyms_other_ deviations	Acronym, Definition, Category	Resolve supported medical and industry acronyms
CROs	Vendor MDM (Databricks)	Dy-namic	us_gmsgq_dev.gms_us_alyt.cro_ deviations	Name, Display Name, description	Resolve known CRO names and add descriptive context
Devices	CTMS (available elsewhere)	Static	us_gmsgq_dev.gms_us_alyt.rd_device_list_ deviations	Name, Alias, Device_Long_name, Device_Type, business_owner	Resolve device-name variants and add type and ownership context
Vendors	Vendor MDM (Databricks)	Dy-namic	us_gmsgq_dev.gms_us_alyt.supplier_list_ deviations	Supplier_Name_RD, Country_RD	Resolve known external supplier names and country context

Entity-Specific Resolution Rules

Clinical Identifiers

Clinical identifiers are resolved from three source paths:

1. Pattern-based extraction from the combined deviation text.
2. Values in `Study_Protocol`, split on semicolons and treated as trusted clinical values after removal of null-like entries.
3. Values in `Program_Number`, split on semicolons and retained only when the leading token matches a supported clinical-ID family.

The implemented identifier families include supported TAK, MLN, SHP, HGT, CCT, DEN, and internal C-number patterns. Both canonical complete identifiers and normalized parent or compound keys may be retained for matching.

Controlled Documents

The pipeline extracts supported SOP, SPEC, MTHD, MTD, PROC, TOOL, FORM, and WI identifiers from free text and `Document_or_Process`. Current, previous, and legacy document-number fields from the reference input are normalized to a common key format.

Acronyms

Acronym extraction follows a known-set approach:

- Candidate uppercase expressions are detected in source text.
- Only candidates present in the normalized acronym reference are retained.
- One- and two-character entries are excluded.
- Values resembling supported clinical or document identifiers are excluded to avoid entity-type collisions.
- Medical and Industry definitions are retained.
- Selected British spelling variants are normalized before definition deduplication.

CROs, Devices, and Vendors

CROs, devices, and vendors are matched against known normalized names using substring matching:

- CRO names shorter than four characters are excluded.
- Device names, aliases, and long names shorter than four characters are excluded.
- Vendor names shorter than four characters are excluded.

- Device business-owner abbreviations are expanded where mappings are configured.
- Matching remains dependent on the coverage and quality of each reference table.

These methods favor interpretable, reference-grounded matches. They do not constitute comprehensive named-entity recognition for all possible organizations, devices, suppliers, or terminology variants.

Output: Text Tiers

The pipeline constructs three source-text representations for every retained event.

Tier	Implementation field	Included source fields	Intended role
Core	<code>source_free_text_core</code>	Event_Title and Event_Description	Concise representation of the primary deviation narrative
Mid	<code>source_free_text_mid</code>	Core fields plus Impact_Assessment, Quality_Final_Assessment, Root_Cause_Category, and Root_Cause_SubCategory	Event-level representation with assessment and root-cause context
Full	<code>source_free_text_full</code>	Mid fields plus aggregated Action_Text	Most detailed representation, including action-record information

The BGE-M3 and AI Search implementation also refers to the full tier as the fine tier. In this document, full describes the text representation and fine index refers to the corresponding implementation object.

Output: Contextualized Embedding

Persistent Enrichment Table

Attribute	Value
Fully qualified table	<code>us_gmsgq_dev.gms_us_alyt.deviation_embed_input</code>
Grain	One row per retained <code>pr_id</code>
Primary purpose	Store source text, deterministic reference context, generated situating context, and final contextual retrieval text
Write behavior	Overwrite through a staging table
Temporary staging tables	<code>deviation_embed_input_staging</code> and <code>deviation_embed_input_stg</code>
Downstream consumers	Embedding generation and retrieval workflow

Output Fields

Field group	Fields	Description
Event identifier	<code>pr_id</code>	String representation of <code>Event_Number</code>
Source text	<code>source_free_text_full</code> , <code>source_free_text_mid</code> , <code>source_free_text_core</code>	Cleaned source text at three levels of detail
Deterministic enrichment	<code>deterministic_context</code>	Resolved and concatenated reference information
Generated enrichment	<code>llm_context</code>	Short model-generated paragraph situating the deviation using source and resolved reference context
Contextual retrieval text	<code>contextual_retrieval_text_full</code> , <code>contextual_retrieval_text_mid</code> , <code>contextual_retrieval_text_core</code>	Final text passed to embeddings and synchronized for keyword retrieval

Each contextual retrieval field is assembled in this order: Resolved deterministic references + LLM-generated situating context + Source free text for the selected tier.

The original source-text tiers remain stored separately so that generated or resolved context does not replace source evidence.

Generated Situating Context

The pipeline calls a Databricks-hosted foundation model through `ai_query`. The configured model identifier is `databricks-gpt-5-4-mini`.

For each non-empty event, the prompt supplies:

- `deterministic_context`
- `source_free_text_full`
- Instructions to produce a short retrieval-oriented description of the entities involved and the subject of the deviation

The configured generation parameters are:

Parameter	Configured value
Maximum generated tokens	300
Temperature	0.0

The generated result is stored as `llm_context` and reused across the full, mid, and core contextual retrieval representations. It is generated context, not source evidence. Reviewers should use the retained source fields and deterministic references when validating its contents.

Deviation Embeddings Table

Attribute	Value
Fully qualified table	<code>us_gmsgq_dev.gms_us_alyt.deviation_embeddings</code>
Grain	One row per <code>pr_id</code>
Primary key	<code>pr_id</code>
Change Data Feed	Enabled
Embedding model	BAAI/bge-m3
Embedding dimension	1,024
Configured maximum sequence length	8,192 tokens
Configured batch size	64
Number of text tiers	Three
Write behavior	Overwrite for the first write chunk; append for subsequent chunks

Embedding Fields

Text tier	Contextual text input	Vector output
Full/fine	<code>contextual_retrieval_text_full</code>	<code>embedding</code>
Mid	<code>contextual_retrieval_text_mid</code>	<code>mid_embedding</code>
Core	<code>contextual_retrieval_text_core</code>	<code>core_embedding</code>

The three text lists are combined into one encoding workload, sorted approximately by token length to reduce unnecessary padding, encoded, restored to their original order, and split back into the three vector columns.

The embedding table also retains:

- `source_free_text_full`
- `source_free_text_mid`
- `source_free_text_core`
- `deterministic_context`
- `llm_context`
- `contextual_retrieval_text_full`
- `contextual_retrieval_text_mid`
- `contextual_retrieval_text_core`

Retaining these fields preserves the relationship between each vector and the text from which the vector was generated.

Output: Normalized Lookup Tables

Entity type	Normalized output table	Normalization and enrichment summary
CLINICAL_ID	us_gmsgq_dev.gms_us_ alyt.ref_clinical_norm	Creates keys from available clinical aliases and adds generic name, modality, therapeutic area, finance grouping, indication, target, and mechanism where available
DOCUMENT	us_gmsgq_dev.gms_us_ alyt.ref_docs_norm	Normalizes current, previous, and legacy document numbers; maps the legacy MTD prefix to MTHD; adds document title
ACRONYM	us_gmsgq_dev.gms_us_ alyt.ref_acronym_norm	Retains supported Medical and Industry definitions, removes short or invalid entries, excludes identifier-like values, and consolidates definitions by acronym
CRO	us_gmsgq_dev.gms_us_ alyt.ref_cro_norm	Uses normalized CRO names as keys and adds name, display name, and description
DEVICE	us_gmsgq_dev.gms_us_ alyt.ref_device_norm	Creates lowercase name, alias, and long-name variants and adds device type and expanded business-owner context
VENDOR	us_gmsgq_dev.gms_us_ alyt.ref_vendors_norm	Uses normalized supplier names as keys and adds supplier name and country
All supported entities	us_gmsgq_dev.gms_us_ alyt.ref_glossary_unified	Unions the six normalized lookup tables and retains one row per key_norm and entity_type

Output: Retrieval

AI Search Endpoint

Attribute	Value
Endpoint name	deviation-retrieval-vs
Endpoint type	STANDARD
Environment	Development
Authentication	Notebook credentials detected by the AI Search client

Delta Sync Indexes

Tier	Index	Vector field	Searchable text field
Full/fine	us_gmsgq_dev.gms_us_ alyt.deviation_emb_fine_idx	embedding	contextual_retrieval_ text_full

Tier	Index	Vector field	Searchable text field
Mid	us_gmsgq_dev.gms_us_alyt.deviation_emb_mid_idx	mid_embedding	contextual_retrieval_text_mid
Core	us_gmsgq_dev.gms_us_alyt.deviation_emb_core_idx	core_embedding	contextual_retrieval_text_core

Each index:

- Uses `deviation_embeddings` as its source table.
- Uses `pr_id` as the primary key.
- Uses a precomputed 1,024-dimensional BGE-M3 vector.
- Uses a triggered Delta Sync pipeline.
- Synchronizes the tier-specific contextual retrieval text for keyword retrieval.
- Does not ask Databricks AI Search to recompute the stored embeddings.

At query time, the retrieval workflow supplies:

- A BGE-M3 query vector for semantic similarity.
- The original query text for keyword matching.
- HYBRID as the query type.

Databricks AI Search returns the combined ranking from vector and keyword retrieval. The notebook returns the top configured results for the selected tier.

Output: Supplementary Table

The supplementary table is the analysis-facing product of the workflow. It is built on top of the formatted TrackWise source and the enrichment tables, collapsing the date- or row-grain source to one row per event and joining the parsed deterministic entities, personnel, location, discovered topics, and root-cause signals to each deviation. It can be joined back to QDN data through the documented event identifier.

Attribute	Value
Fully qualified table	us_gmsgq_dev.gms_us_alyt.deviation_supplementary
Grain	One row per event (<code>pr_id</code> , equal to <code>Event_Number</code>)
Primary purpose	Join enriched entities, personnel, location, topics, and root-cause signals back to each deviation event for analysis and review
Key source	us_gmsgq_dev.gms_us_mart.tw_deviation_data_formatted_rdq (event key, personnel, and location text)
Enrichment source	us_gmsgq_dev.gms_us_alyt.deviation_embeddings (source text, contextual retrieval text, and deterministic context)
Additional sources	deviation_topic_assignments (zero-shot BERTopic topics) and deviation_root_cause (LLM root cause)

Attribute	Value
Write behavior	Overwrite with schema overwrite enabled
Join key	pr_id, cast to string, left-joined across all contributing tables

Output: Supplementary Fields

Field	Fields	Description
Source text group	source_free_text_full Event_Number	Cleaned deviation free text spanning all Event key restored from pr_id; one row per event
Contextual retrieval text	contextual_retrieval_text_full	Enriched text (deterministic references + generated situating context + source text)
Deterministic enrichment	deterministic_context	Raw resolved-references blob retained for reference
Parsed entities	clinical_ids, documents, acronyms, cros, devices, vendors	Typed arrays parsed from deterministic_context, deduplicated in first-seen order and preferring the resolved canonical name
Personnel	personnel	Distinct non-null names collected from Event_Owner, QA_Contact, Reported_by, and Quality_Approver_CAPA_Final
Local Operating Company LOC Array of inferred countries from name, agency, ccTLD, and adjacent-code signals; [LOC] when only the structured operational-unit flag is present, otherwise empty		
Topics	primary_topic, secondary_topic, tertiary_topic	Zero-shot BERTopic (BGE-M3) topic labels; null when the topic table is unavailable
Root cause	llm_rc_primary, llm_rc_secondary, llm_rc_tertiary	LLM-judged primary root cause (Ishikawa 6M) and first two contributing factors; null when the root-cause table is unavailable

The parsed entity, personnel, and LOC fields are stored as string arrays. The topic and root-cause fields are populated by optional upstream tables; when those tables are absent, the pipeline writes the columns as nulls so that the supplementary schema remains stable. The full source and contextual retrieval text are retained alongside the extracted fields so that every derived value can be traced back to its supporting evidence.

Input-to-Output Lineage

Input	Principal transformation	Persistent output	Downstream use
Clinicalid_ deviations	Alias parsing, clinical-ID normalization, contextual-field assembly	ref_clinical_ norm	Unified glossary and deterministic enrichment
documents_number_ deviations	Current, previous, and legacy document normalization	ref_docs_norm	Unified glossary and deterministic enrichment
Acronyms_other_ deviations	Filtering, normalization, definition consolidation	ref_acronym_ norm	Acronym extraction, glossary enrichment, and selected abbreviation expansion
cro_deviations	CRO-name filtering and normalization	ref_cro_norm	Known-set CRO extraction and enrichment
rd_device_list_ deviations	Name, alias, and long-name normalization	ref_device_ norm	Known-set device extraction and enrichment
supplier_list_ deviations	Supplier-name normalization	ref_vendors_ norm	Known-set vendor extraction and enrichment
Six normalized references	Schema alignment, union, and deduplication	ref_glossary_ unified	Single reference join for all supported entity types
tw_deviation_data_ formatted_rdq	Event aggregation, text cleaning, language filtering, text-tier construction, and entity extraction	deviation_ embed_input	Context generation and embedding
deviation_embed_ input	BGE-M3 dense encoding of three contextual text tiers	deviation_ embeddings	Vector and hybrid retrieval
deviation_ embeddings	Triggered Delta Sync indexing	Three AI Search indexes	Core, mid, and full/fine retrieval

Output: Temporary and Diagnostic

The following objects or results support safe execution and validation but are not canonical downstream data products.

Artifact	Purpose	Persistence
deviation_embed_ input_staging	Materializes the event-level embedding input before replacing the	Temporary; dropped after the write canonical table
deviation_embed_ input_stg	Materializes model-generated context before replacing the	Temporary; dropped after the write canonical table
Entity extraction coverage	Counts events with at least one extracted key by entity type	Notebook diagnostic
Reference join coverage	Compares extracted keys with keys matched to the unified glossary	Notebook diagnostic

Artifact	Purpose	Persistence
Unmatched-key report	Lists frequently extracted keys without a reference match Notebook diagnostic	
LLM-context coverage	Counts populated and empty <code>llm_context</code> values	Notebook diagnostic
Embedding integrity checks	Checks row counts, distinct <code>pr_id</code> values, vector dimensions, Notebook diagnostic and context coverage	
Index-status checks	Reports index state and indexed-row counts Notebook diagnostic	

Technical Dependencies

The pipeline was developed and tested on Databricks Serverless compute. Because the `%pip install` cells do not pin exact versions, the versions below were captured from the running environment for reproducibility.

Runtime environment

Item	Value
Compute type	Databricks Serverless
Serverless environment ver.	[capture] (Notebook → Environment panel)
Python	3.12.3
GPU / CUDA	Required for FP16 BGE-M3; CUDA 12.9 (torch 2.9.0+cu129)

Python packages (captured versions)

Package	Version	Role
FlagEmbedding	1.4.0	BGE-M3 model loading and dense vector generation
torch	2.9.0+cu129	GPU tensor backend (CUDA 12.9) for BGE-M3
transformers	4.57.6	Model and tokenizer dependency
bertopic	0.17.4	Zero-shot topic modeling
umap-learn	0.5.9.post2	Dimensionality reduction for BERTopic
hdbscan	0.8.44	Density-based clustering for BERTopic
scikit-learn	1.6.1	CountVectorizer (c-TF-IDF) and cluster metrics
databricks-ai-search	0.78	Endpoint and Delta Sync index management and querying
openai	3.0.0	Client for the LLM classification endpoint
langdetect	1.0.9	Reproducible English-language filtering
numpy	2.1.3	Vector-order restoration and array handling
pandas	2.2.3	Driver-side transfer for embedding generation
PySpark	(Serverless)	Source aggregation, joins, and Delta writes

Package	Version	Role
---------	---------	------

Model and endpoint configuration

Item	Value
Embedding model	BAAI/bge-m3
Embedding model revision (commit)	[capture] (Hugging Face snapshot hash)
Embedding dimension	1,024
Max sequence length	8,192 tokens
Situating-context FM endpoint	databricks-gpt-5-4-mini (ai_query)
Situating-context params	temperature=0.0, max_tokens=300
Topic-classifier FM endpoint	databricks-gpt-5-4-mini (OpenAI client)
Topic-classifier params	temperature=0.0, max_tokens=350
Vector Search endpoint	deviation-retrieval-vs (type STANDARD)

Running the Code

The workflow is a set of Databricks notebooks that must be run in **dependency order**: each writes Delta tables or vector indexes that later notebooks read. Run the cells top-to-bottom within each notebook. `build_reference_glossary` is the foundation and must run first; the reporting notebook runs last.

1. `build_reference_glossary.ipynb`. Normalizes the six reference datasets into `ref_glossary_unified`, compresses the deviation source to event grain, resolves entities, generates the LLM situating context, builds the three-tier BGE-M3 embeddings (`deviation_embeddings`), and creates the AI Search endpoint and the three Delta Sync indexes. Everything downstream depends on its outputs.
2. `topic_modeling.ipynb`. Reads `deviation_embeddings` and assigns interpretable topics with zero-shot BERTopic, writing `deviation_topics` and the primary/secondary/tertiary labels in `deviation_topic_assignments`.
3. `root_cause.ipynb`. Reads `deviation_embeddings` plus the human `Root_Cause_Category` from the source mart and writes the LLM Ishikawa (6M) root-cause judgement per event to `deviation_root_cause`. It is independent of `topic_modeling`—once step 1 is done, steps 2 and 3 may run in either order.
4. `create_supplementary_table.ipynb`. Joins the resolved entities, personnel, location, topics (step 2), and root cause (step 3) into the analysis-facing `deviation_supplementary`, one row per event. Requires steps 1–3.

5. `unblinding_semantic_search.ipynb`. A read-only consumer of the retrieval stack: builds thematic cohorts (inadvertent unblinding, protocol deviations) by unioning a native-HYBRID semantic search with a keyword regex, and writes `unblinding_cohort`. Requires the indexes from step 1; independent of steps 2–4.
6. `deviation_statistics.ipynb` — run last. The reporting layer. Reads the supplementary table, topics, root causes, and the unblinding cohort to produce summary statistics, figures, SME review worklists, and the curated `deviation_triage` export. Requires steps 1–5.

Justification of Methodology

Retrieval-augmented generation splits documents into short chunks, embeds each as a dense vector, and retrieves the chunks most similar to a query [1]. Chunking strips away surrounding context—an acute problem for quality text written in Takeda-internal syntax. A deviation description may reference a study, compound, controlled document, device, or vendor only by an abbreviation or internal identifier that means little on its own; as of August 2026, 96.7% of entries mention at least one such entity. For example, `pr_id=5225817` reads, “A US site participating in the TAK-999-3002 study attempted to randomize a participant...”—but without context, neither a retriever nor a topic model can tell whether TAK-999-3002 is a study, compound, document, device, or vendor.

Contextual Retrieval addresses this by prepending a short, chunk-specific description to each unit of text before it is embedded and indexed, producing Contextual Embeddings and Contextual BM25; because manual annotation is infeasible at scale, a language model generates the description from surrounding context [1]. The pipeline mirrors this: for each deviation, a foundation model generates a brief situating description (`llm_context`) from the resolved reference context and source narrative, which is combined with the deterministic reference context and original text to form the contextual retrieval text that is embedded and indexed. This context is stored separately, so generated text supplements rather than replaces the source evidence.

A purely semantic index is insufficient, because dense embeddings capture meaning but miss exact matches, whereas BM25 recovers them. This is decisive for deviation records dense with protocol numbers, document numbers, and internal codes that must match literally, where even the final digit distinguishes vastly different scenarios [1]. A controlled scaling study over 511,959 documents (about 601M tokens) reinforces this: holding questions and evidence fixed while growing the corpus across 28 tiers, BM25 overtakes an LLM-driven agent at roughly 10M tokens and leads at every larger tier, approaching a 20-point margin at full scale while anchoring the low-cost end of the accuracy–cost frontier [2]. The advantage holds because enterprise questions hinge on precise lexical anchors while distractors are semantically similar but wrong; it persists under paraphrase, and BM25’s query cost stays nearly constant (about 5.8K tokens) as the corpus grows [2].

This justifies two design choices: situating context is generated and prepended before indexing, so every embedding and keyword entry carries the context needed for reliable retrieval; and retrieval runs in **HYBRID** mode, fusing dense similarity with BM25 over the same contextualized text. Combining semantic and lexical signals outperforms embeddings alone [1], and the scaling evidence marks lexical ranking as the strongest scalable default for enterprise corpora [2], making BM25 essential for the exact-identifier queries that characterize deviation analysis.

Topic Modeling and its Application to Clinical and Quality Data

Topic modeling is an unsupervised technique that discovers latent themes in a document collection without predefined labels [3]. It assumes each document is a mixture of topics and each topic a distribution over words, so a topic’s highest-probability words convey what it is about; because it summarizes at the document level, it is well suited to exploring large volumes of unstructured text [3].

Latent Dirichlet Allocation (LDA) [4] makes this precise as a three-level generative model. For a corpus of M documents over a vocabulary of size V with k topics, each topic’s word probabilities are fixed by a $k \times V$ matrix β with $\beta_{ij} = p(w^j = 1 \mid z^i = 1)$. A document of N words is generated by drawing topic proportions $\theta \sim \text{Dir}(\alpha)$ on the $(k - 1)$ -simplex, then, for each word, a topic $z_n \sim \text{Multinomial}(\theta)$ and a word $w_n \sim p(w_n \mid z_n, \beta)$ [4], giving

$$p(\theta, \mathbf{z}, \mathbf{w} \mid \alpha, \beta) = p(\theta \mid \alpha) \prod_{n=1}^N p(z_n \mid \theta) p(w_n \mid z_n, \beta)$$

The document likelihood follows by integrating over θ and summing over $\mathbf{z}$; the resulting posterior is intractable and is approximated by variational EM [4]. Model quality is typically judged by topic coherence and topic diversity, and common challenges include choosing the topic count, instability across runs, and extracting coherent themes from short, sparse texts [3].

BERTopic reframes topic modeling as clustering over contextual embeddings, and separates document clustering from topic representation so language models can be swapped in and either stage tuned independently [5]. Each document is embedded with a pre-trained transformer; the embeddings are reduced with UMAP and clustered with HDBSCAN, a density-based method that finds clusters of varying density and treats unrelated documents as outliers rather than forcing them into a topic [5]. Topic words are then derived per cluster with a class-based TF-IDF, generalizing $W_{t,d} = \text{tf}_{t,d} \log(N/\text{df}_t)$ to $W_{t,c} = \text{tf}_{t,c} \log(1 + A/\text{tf}_t)$, where A is the average number of words per class; the topic count need not be fixed in advance [5].

BERTopic suits deviation narratives for three reasons. It works on contextual embeddings rather than bag-of-words counts, capturing semantic relationships LDA and NMF discard, and it reuses the same BGE-M3 representation the pipeline already computes, so topics are derived from enriched, entity-resolved text [5]. It does not require a preset topic count: HDBSCAN

discovers clusters of varying density and isolates genuine one-off events as outliers, fitting a heterogeneous deviation corpus better than LDA’s, which, due to its a-priori count, has the tendency to assign one dominant topic [5]. Its main limitation, one topic per document, is mitigated by retaining primary, secondary, and tertiary labels from HDBSCAN’s soft-cluster probabilities, the author’s own suggested workaround, while the interpretable c-TF-IDF topic words remain reviewable by SMEs for downstream root-cause and trend analysis [5].

Topic modeling has been applied across information retrieval, text classification, recommendation, and bioinformatics [3]. Its relevance to quality data is direct: Sánchez-Franco et al. applied BERTopic to quality-management abstracts (1980–2020), recovering 48 topic-clusters and 13 meta-topics and tracing their evolution [3], the same mixture-of-topics view that treats a deviation as a blend of recurring operational themes rather than a single rigid category. But quality free text is dense with context-dependent short forms: “MS” can mean morphine sulfate, multiple sclerosis, or mass spectrometry, and the same form can differ across sites [6]. Unresolved, this ambiguity fragments the vocabulary a topic model sees and degrades coherence [3], which is precisely why the pipeline resolves references and normalizes acronyms before topics are assigned, yielding cleaner, more interpretable themes.

The same ambiguity undermines retrieval. The mechanism that gathers related deviations for topic and root-cause analysis. A chunk mentioning only a bare code or acronym carries too little context to match a query reliably, the same context-loss problem Contextual Retrieval solves by prepending a situating description before embedding and indexing [1]. As detailed in the earlier section, grounding each record therefore does double duty: it disambiguates the vocabulary the topic model sees and enriches the text the retriever indexes, which is why the pipeline pairs contextual enrichment with hybrid dense-plus-BM25 retrieval, supporting both the themes surfaced by topic modeling and the exact-identifier queries characteristic of deviation analysis [2].

Methods: Topic Modeling and Root-Cause Assignment

The pipeline assigns interpretable topics and root-cause categories to each deviation event using three complementary strategies over the same contextualized BGE-M3 representation: (1) zero-shot BERTopic against a fixed quality taxonomy, (2) an unsupervised full- vs. core-tier BERTopic comparison, and (3) an LLM-based multi-label classifier. All three consume the enriched text produced by the contextual-retrieval stage, so entities and acronyms are resolved before topics are assigned.

Topic assignment reads from `deviation_embeddings`, the only table holding both the pre-computed BGE-M3 vectors and the source text. Two design decisions follow directly from BERTopic’s architecture, which deliberately separates document clustering from topic representation [5]. First, clustering uses the `embedding` column, BGE-M3 encoded over the enriched `contextual_retrieval_text_full`, so documents are grouped in a space that already carries resolved-reference and situating context, directly addressing the context-loss and short-form

ambiguity that would otherwise fragment topics [1], [6]. Second, the class-based TF-IDF (c-TF-IDF) labeling step uses the clean `source_free_text_full`, stripped of its `Column_Name`: scaffolding, so topic keywords reflect deviation content rather than schema tokens. Because the embeddings are precomputed, BERTopic skips re-encoding and only encodes the short candidate-label strings.

The first strategy defines 18 candidate topics from GxP quality vocabulary (inadvertent unblinding, GMP manufacturing deviation, stability OOS). Each label is encoded with BGE-M3 into the same 1024-dimensional space as the documents, and events are matched by cosine similarity above a 0.5 threshold; documents matching no label are clustered as emergent topics via UMAP + HDBSCAN. This is a faithful instantiation of the BERTopic procedure: transformer embeddings, UMAP dimensionality reduction—chosen because distance measures degrade in high dimensions—and HDBSCAN density-based clustering that models unrelated documents as outliers rather than forcing them into a topic [5]. Topic words are then generated per cluster with c-TF-IDF, following the paper’s generalization of TF-IDF from documents to concatenated document “classes” [5].

A deviation often spans several themes (a manufacturing deviation, an overdue CAPA, and a documentation gap). This is precisely BERTopic’s acknowledged limitation—it assigns one dominant topic per document—and the mitigation implemented here is the author’s own suggested workaround: using HDBSCAN’s soft-cluster probabilities to recover a topic distribution per document [5]. Concretely, `approximate_distribution()` computes topic relevance over a sliding window of passages, and topics above a 0.1 threshold are retained as primary, secondary, and tertiary labels, written to `deviation_topic_assignments`. This is the same mixture-of-topics view that motivates LDA, where a document is modeled as a mixture over latent topics rather than a single category [4].

To test the value of contextual enrichment, the same BERTopic configuration (identical UMAP/HDBSCAN parameters and c-TF-IDF vocabulary) is fit twice: once on `embedding` (BGE-M3 over the full contextualized text) and once on `core_embedding` (BGE-M3 over the core text only). Agreement between the two partitions is quantified with Adjusted Rand Index and Normalized Mutual Information, and quality is compared via outlier rate and topic diversity, which is the distinctiveness metric standard in the topic-modeling literature [3], [5]. Holding the clustering algorithm fixed isolates the effect of how much context was embedded, testing the central claim that prepending situating context improves the representation [1].

The third strategy is a per-document, multi-label LLM classifier (`databricks-gpt-5-4-mini`) run over `contextual_retrieval_text_full`. A hybrid prompt prefers a fixed 21-topic taxonomy but may invent a concise novel topic when nothing fits, and additionally returns a root-cause category (an Ishikawa-style controlled vocabulary), a regulatory-impact category, a quoted evidence span, and a calibrated confidence score. Grounding guardrails instruct the model to use “Not determinable” rather than guess, keeping generated labels tied to textual evidence. Returned labels are snapped to canonical spellings; unmatched labels are tracked as taxonomy-expansion candidates. Running the classifier on the enriched text rather than raw

narrative is essential: without resolved references, an identifier such as **TAK-999-3002** carries too little standalone meaning for either a topic model or an LLM to interpret reliably [6].

Finally, zero-shot BERTopic and LLM labels are compared per event across primary, secondary, and tertiary assignments. The two methods are complementary: BERTopic provides unsupervised, corpus-grounded clusters with reproducible c-TF-IDF keywords, while the LLM provides evidence-anchored labels and root-cause reasoning. Agreement raises confidence in an assignment; divergence flags events for subject-matter-expert review, supporting the downstream root-cause and trend analysis the supplementary table is built for.

Feature: Semantic Search

The unblinding cohort builder shows how to query the retrieval stack to assemble a high-recall set of related deviations for a business question—here, inadvertent unblinding. It is a read-only consumer of the indexes built upstream; it retrieves and profiles events but never reprocesses raw data. The same pattern generalizes to any thematic cohort (protocol deviations, data-integrity events, etc.) by swapping the queries and keyword pattern.

How to run it

1. Provision a GPU. BGE-M3 is loaded only to encode the handful of query phrasings; the keyword path needs no model, so a keyword-only run can skip the GPU entirely.
2. Define the query set. Edit `UNBLINDING_QUERIES` with several phrasings that capture how the theme surfaces in practice (“code break performed without authorization,” “IP kit revealed treatment allocation”). More phrasings raise recall.
3. Define the keyword pattern. Edit the `KW` regex with literal terms and spelling variants (`unblind`, `un-blind`, “blind was broken”).
4. Run the retriever. `CohortRetriever` encapsulates the AI Search client, the BGE-M3 encoder, and the index config; `semantic_ids()` and `keyword_ids()` return the two `pr_id` sets, which are unioned into the cohort.
5. Inspect the union counts. The printout of *both signals / semantic-only / keyword-only* is the first quality check—if semantic-only dominates, lower `SEM_PER_QUERY` or sharpen the phrasings; if it is near-empty, raise `SEM_PER_QUERY` or add phrasings.

The cohort is joined back to the source mart for the rich fields (`Root_Cause_Category`, `Impact_Assessment`, `Action_Text`, program/protocol) and written to `unblinding_cohort` for downstream reporting.

The semantic leg runs each phrasing as a native AI Search **HYBRID** query, fusing dense BGE-M3 similarity with Okapi BM25 over the same contextualized text via Reciprocal Rank Fusion; the separate keyword regex is unioned on top as a guaranteed exact-term precision anchor. Because RRF-fused scores are not cosine-comparable, recall is governed by how many top hits each phrasing contributes rather than a similarity cutoff:

Parameter	Effect	When to adjust
SEM_PER_QUERY (default 50)	Top-N HYBRID hits pulled per phrasing	Raise for higher recall; lower to keep the cohort tight

Tune it by inspecting borderline hits, not by guessing: pull a sample of semantic-only matches and check whether they are genuine unblinding events.

The builder profiles the cohort along the dimensions the business questions ask for— cohort size as a share of all deviations, impacted programs and protocols, recurrence (programs with 2 events), root-cause categories and sub-categories, CAPA/**Action_Text** for review, and keyword-flagged safety and data-integrity impact. The provenance pie and root-cause/program bars summarize this visually.

Before quoting the cohort percentage externally, spot-check the semantic-only hits: keyword-only and both-signal events are high-confidence, but semantic-only matches are the ones most likely to include false positives. The impact figures are keyword flags over inconsistent free text and are a screening signal, not a validated impact determination—the **Action_Text** and impact fields are surfaced for SME review precisely because a full answer (“are CAPAs preventing recurrence?”) needs CAPA-status data the source does not contain.

This was used to great success for topics about unblinding and protocol deviations.

Feature: Configured Heuristics and Assumptions

The reporting layer applies several hard-coded rules, thresholds, and vocabularies. These are analyst- and SME-defined configuration, not model outputs: they encode how the business (per the CPMQ business-question set) wants deviations grouped, flagged, and prioritized. They are documented here so reviewers can audit and adjust them.

Business-category taxonomy

`topic_category()` maps each BERTopic label to one of nine high-level categories (Clinical / GCP, Data Integrity & IT, Manufacturing / GMP, Equipment & Facilities, Quality Control, Quality Systems, Supply Chain, Safety & PV, Other / Emergent) using keyword rules. This rollup answers the “where is delivery breaking down” and “which issue categories recur across sources” questions by collapsing dozens of fine-grained topics into a stable, reviewable set of themes. The mapping is keyword-based and deliberately robust to label truncation; unmatched labels fall through to Other / Emergent rather than being force-fit.

Severity / regulatory-impact flags

SEVERITY_PATTERNS applies four case-insensitive regex flags over `Impact_Assessment` + `Quality_Final_Assessment`:

Flag	Triggers on (keywords)
Patient safety	safety, subject, patient, harm, injur, adverse
Data integrity	data integrity, GCP/GxP, validit, falsif, ALCOA
Regulatory	regulator, FDA/EMA/ICH, inspection, reportab, recall
Product quality	batch, contaminat, OOS, sterilit

These let reviewers triage by risk rather than by volume, directly supporting the safety- and data-integrity-impact requirements of Q1 and the inspection-exposure questions. The patterns are keyword heuristics over free text and are intended as a first-pass screen, not a compliance determination.

Repeat-offender and recurrence thresholds

- `recur >= 2` defines a “repeat-offender” program or protocol (2 deviations).
- Recurring `topic root-cause` signatures are the most frequent (topic, root cause) pairs.

Both operationalize the many CPMQ questions about systemic vs. isolated issues and recurrence across programs / CROs / vendors. A threshold of 2 is the minimal definition of “recurring”; SMEs may raise it to focus on the most systemic offenders.

Automation-value assumption

`MIN_PER_EVENT = 20` is an SME-provided estimate of the analyst minutes required to manually triage and root-cause one deviation. It drives the “analyst-hours saved” tiles and is the single assumption behind the effort-savings figures. It is clearly an estimate; the hours-saved numbers scale linearly with it, so it should be revised to each team’s actual baseline before being quoted externally.

Validation thresholds

- Confidence 80 (or textual “high”) marks an event as high-confidence auto-accept; everything else routes to analyst review.
- `review_priority` sums the review drivers (outlier topic + low confidence + model human root cause + severity score), so the queue sorts worst-first.
- `review_reasons` records why each event was flagged.

This is the human-in-the-loop validation funnel. The pipeline auto-accepts only where it is confident and agrees with the human record, and escalates disagreements and high-severity events to support the CAPA-effectiveness and “predictive vs. reactive” questions by making model trust explicit and auditable.

Root-cause label reconciliation

LLM_TO_GT maps the LLM’s Ishikawa 6M labels to the reviewer-entered `Root_Cause_Category` naming before comparison, so the agreement confusion matrix compares like with like. This underpins the model-vs-human agreement metric that establishes whether the automated root-cause call can be trusted.

Semantic cohorts

The unblinding cohort is consumed pre-built from `unblinding_cohort` (hybrid vector + keyword retrieval) to serve the unblinding-focused questions. The protocol-deviation cohort is defined as any event whose primary, secondary, or tertiary topic mentions “protocol,” reusing the BERTopic signal rather than a new text pass, to serve the protocol-deviation questions.

Path to Production

The out-of-scope capabilities are deferred to a separate productionization and governance effort. The steps below are Databricks-native and require no redesign of the proof of concept—only promoting it under managed orchestration, governance, and monitoring.

Because the pipeline currently reprocesses the full corpus on each run, the transformation stages should be migrated to Lakeflow Declarative Pipelines (formerly Delta Live Tables) so that processing becomes incremental [7], [8]. Each stage would become a flow writing to a streaming table, with source updates applied through AUTO CDC rather than plain appends [8]—relying on Change Data Feed, which is already enabled on `deviation_embeddings`—so that only changed events are re-enriched and re-embedded.

The three Delta Sync indexes are similarly refreshed on a triggered pipeline today; switching the latency-sensitive tiers to a continuous pipeline would keep them current as new data arrives [9], [10]. Since they read precomputed vectors from a Change-Data-Feed-enabled table, those updates are incremental and require no re-encoding of unchanged rows [10], though the higher cost of continuous operation means it is best reserved for the fine tier while the mid and core tiers stay triggered.

In production, the notebook diagnostics would be replaced by Databricks Lakehouse Monitoring [11], [12], pairing a time-series monitor over the event tables—tracking null rates, entity coverage,

and `llm_context` population—with an inference monitor that compares the topic and root-cause outputs against the human `Root_Cause_Category` baseline to follow drift and agreement over time, complemented by system tables for pipeline-health and cost observability [13]. Sustained drift then serves as the signal to re-fit BERTopic, refresh the candidate topics, or revise the LLM prompt.

Finally, governance would center on Unity Catalog: the derived products would be promoted from the development catalog into a dev/staging/prod structure [12], [14], protected with role-based access alongside row filters and column masks on sensitive text [12], and guarded at the record level with pipeline expectations [7], [8].

Future Work

The current pipeline uses BGE-M3 in dense mode only. BGE-M3 natively supports multi-vector (ColBERT-style) and sparse representations alongside its dense output, so a natural extension is to exploit late-interaction retrieval, that is, scoring queries against token-level vectors rather than a single pooled embedding.

Acknowledgements

I am grateful to have had the opportunity to work on this project, and I hope it will be useful to the people who rely on deviation data for their work.

This work was made possible through collaboration across R&D Quality Compliance, Systems and Data, and DD&T. I am grateful to my manager, Saeed Babaeizadeh, and to Aayushi Prasad, my Tomodachi. I thank them for sharing their expertise on TrackWise, deviations, and deterministic rules; Calvin Neprud, Karthik Suresh, Michael Prokle, and Bob Sorabella for helping me locate the reference tables for MDM. I owe particular thanks to Katie Evenson and Calvin. Their guidance shaped this project in lasting ways. Beyond their feedback on seeding topics, business logic, and schema planning, they helped me understand what genuinely useful outputs look like and how to frame them for the people who would rely on them, and Katie's encouragement and tips on presenting the work gave me both direction and confidence, and this project is stronger for it. My thanks also to Jared Wiemann and Kanrong Yu for their support of this work.

I am grateful for the many conversations along the way with Mingxi Song, Kevin Cheng, Minnan Xu, Neda Hassanpour, Nellie Holger, and everyone I met across MA Bio Ops, DD&T, and over lunch. Finally, I would especially like to thank Xuemei Li and Yunliang Cai for their mentorship. They gave me a great deal of perspective.

This project was, for me, far more than a technical deliverable. I came away with a more independent (agentic, if you will) problem-solving workflow and a clearer sense of how to set goals and work toward them and learned what it means to build within data-governance constraints. I was never drawn to NLP, yet somehow now, I cannot stop thinking about it. Retrieval-system design and the Databricks documentation page plague my nightly visions. I realized I was interested in something in the domain of NLP after all: Optimal guarantees for topic models with an unknown number of topics using NMF.

I will head into my final year of undergraduate study at the University of Toronto. My primary interest remains in applying machine learning and multivariate statistics to medicine and biology; functional analysis and representation learning. If you would like to continue the conversation, please feel free to reach out to me on LinkedIn.

References

Generative AI tools were used to assist with code, drafting, and formatting. All final decisions and content are my own. Quarto was used to render this document. The references were generated with Zotero and the Better BibTeX plugin and formatted with the apa CSL style.

- [1] D. Ford, “Introducing contextual retrieval.” Accessed: Aug. 13, 2026. [Online]. Available: <https://www.anthropic.com/engineering/contextual-retrieval>
- [2] P. Wang *et al.*, “BM25 wins at scale: A scaling study of retrieval-augmented generation paradigms.” 2026. Available: <https://arxiv.org/abs/2607.26497>
- [3] M. Hankar, M. Kasri, and A. Beni-Hssane, “A comprehensive overview of topic modeling: Techniques, applications and challenges,” *Neurocomputing*, vol. 628, p. 129638, 2025, doi: 10.1016/j.neucom.2025.129638.
- [4] D. M. Blei, A. Y. Ng, and M. I. Jordan, “Latent dirichlet allocation,” *Journal of Machine Learning Research*, vol. 3, pp. 993–1022, 2003.
- [5] M. Grootendorst, “BERTopic: Neural topic modeling with a class-based TF-IDF procedure,” *arXiv preprint arXiv:2203.05794*, 2022, Available: <https://arxiv.org/abs/2203.05794>
- [6] A. Kugic *et al.*, “Processing of short-form content in clinical narratives: Systematic scoping review,” *Journal of Medical Internet Research*, vol. 26, p. e57852, 2024, doi: 10.2196/57852.
- [7] “Spark declarative pipelines (lakeflow pipelines).” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/aws/en/ldp/>
- [8] “Streaming tables.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/aws/en/ldp/concepts/streaming-tables>
- [9] “Create an index — vector search indexes API.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/api/workspace/vectorsearchindexes/createindex>
- [10] “Vector search indexes API reference.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/api/workspace/vectorsearchindexes>
- [11] “Monitor data and AI assets with lakehouse monitoring.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/aws/en/lakehouse-monitoring>
- [12] “Data governance with unity catalog.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/aws/en/data-governance/>
- [13] “Design observability strategy.” Accessed: Aug. 14, 2026. [Online]. Available: <https://learn.microsoft.com/en-us/azure/databricks/lakehouse-architecture/deployment-guide/observability>
- [14] “Best practices for data and AI governance.” Accessed: Aug. 14, 2026. [Online]. Available: <https://docs.databricks.com/gcp/en/lakehouse-architecture/data-governance/best-practices>
- [15] J. Liu *et al.*, “LLM-guided semantic-aware clustering for topic modeling,” in *Proceedings of the 63rd annual meeting of the association for computational linguistics (volume 1: Long papers)*, Association for Computational Linguistics, 2025, pp. 18420–18435.

- [16] Y. Liu *et al.*, “Learning decomposed contextual token representations from pretrained and collaborative signals for generative recommendation,” in *Proceedings of the 49th international ACM SIGIR conference on research and development in information retrieval (SIGIR '26)*, ACM, 2026, pp. 1–11. doi: [10.1145/3805712.3809578](https://doi.org/10.1145/3805712.3809578).
- [17] M. A. Nguyen, K. Sai, and M. Nguyen, “Tracing the evolution of word embedding techniques in natural language processing,” *arXiv preprint arXiv:2603.13271*, 2026, Available: <https://arxiv.org/abs/2603.13271>